

CSI4107
Assignment 2

Junhan Liu 7228243
Yinggin Guan 7285769

Introduction

Our choice of programming language to extract features from the twitter collection is Python. Our script to generate the according data representation that led to the best result for all three classifiers is located in folder ‘python script’. Generally, the script is able to process each twitter from the collection to get the according features including: positive sentiment score, negative sentiment score, neutral score, compound score, positive/negative words count, contains question mark or not, polarity score of a twitter. Other functionalities for part 1, which uses bag of words representation, includes removing stop words, rare words, emojis, hyperlinks, and all special characters. Detailed explanation of what features are chosen for an experiment is included in each section below. Works were distributed equally. Me and my partner were both programming the script and exploring existing methods to improve our experiment results.

How To Run

Location of script: ‘python script’ folder

Script to run: generate.py

Before you run, you should have python 3.6 installed on your PC. Please do not use Python idle to run the script, under some platforms, it might give you python building error and might result in not able to run script properly.

To be able to get the best experiment result, we applied an external python library called NLTK. The API reference is listed in a REMDME.md file in the according experiment folder. Thus, the necessary external packages might be downloaded before running. Packages you might need to have are: nltk, SentimentIntensityAnalyzer, tweepy, and TextBlob. Once you have downloaded Python, you can use >pip install ‘whatever you need’ to install the missing packages.

Part 1

For part 1, “semeval_twitter_data.arff” from the original package is used for weka preprocessing. StringToWordVector filter and Remove duplicate is used to make the string to be extracted to numeric type structure for better performance of classification. As showed in Figure 1.1 to 1.2, there are 2369 attributes with unnecessary features (attributes). After that, the 3 classifiers (SVM, NB and DT) is used and the results are showed in Figure 1.3 to 1.5.

Figure 1.1 StringToWordVector

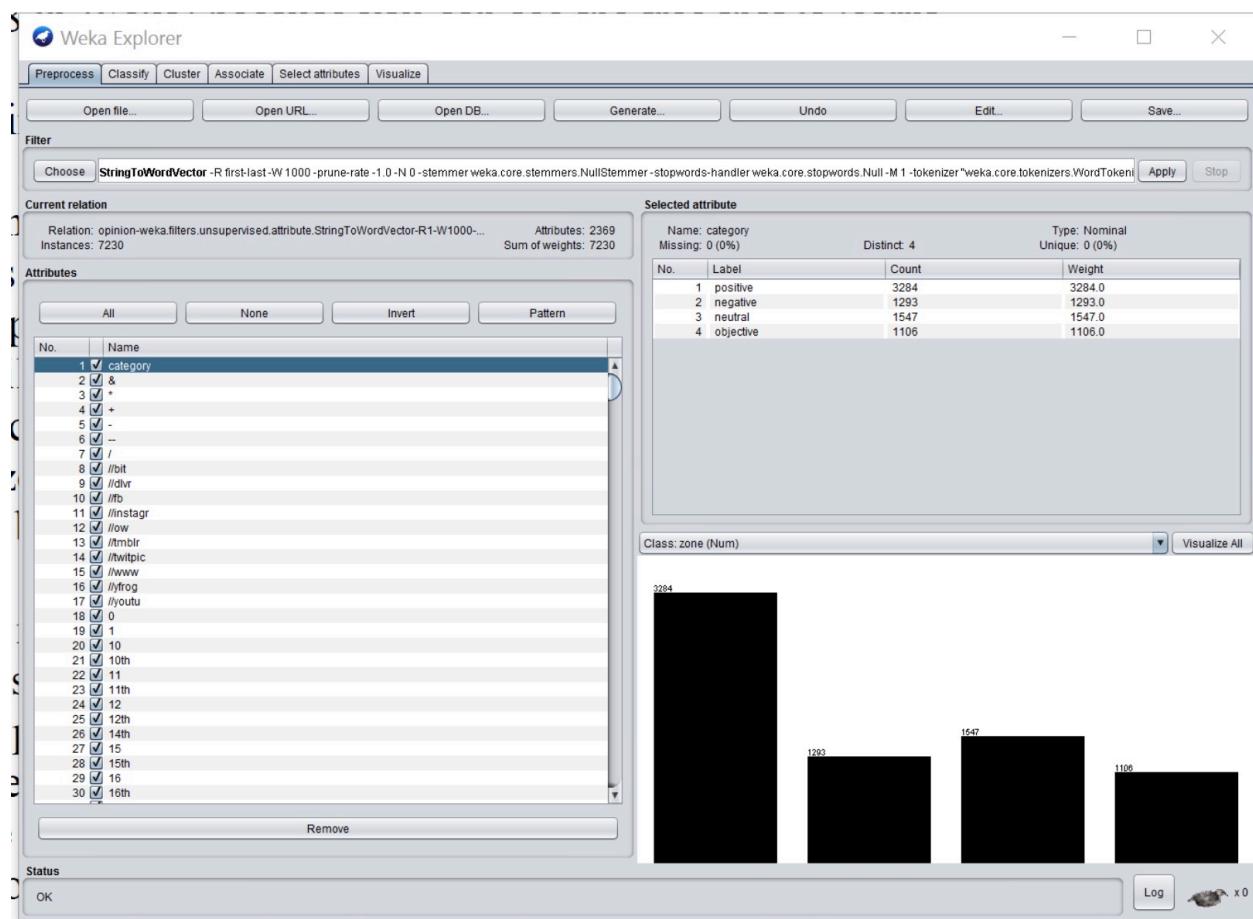


Figure 1.2 RemoveDuplicate

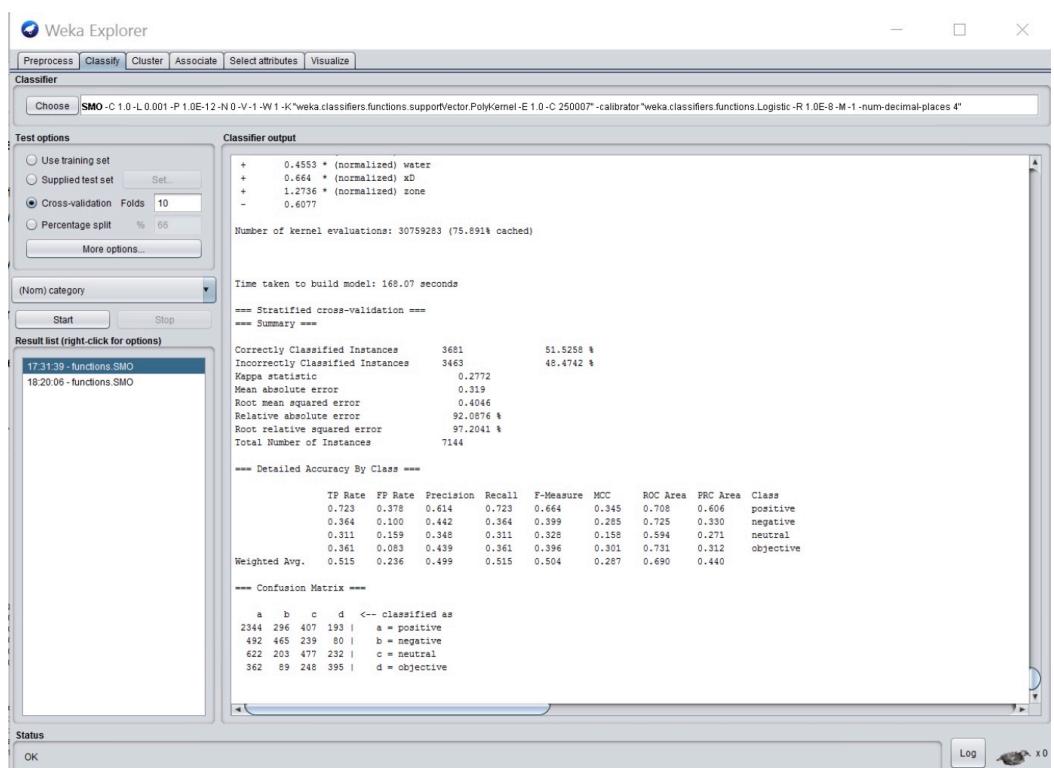
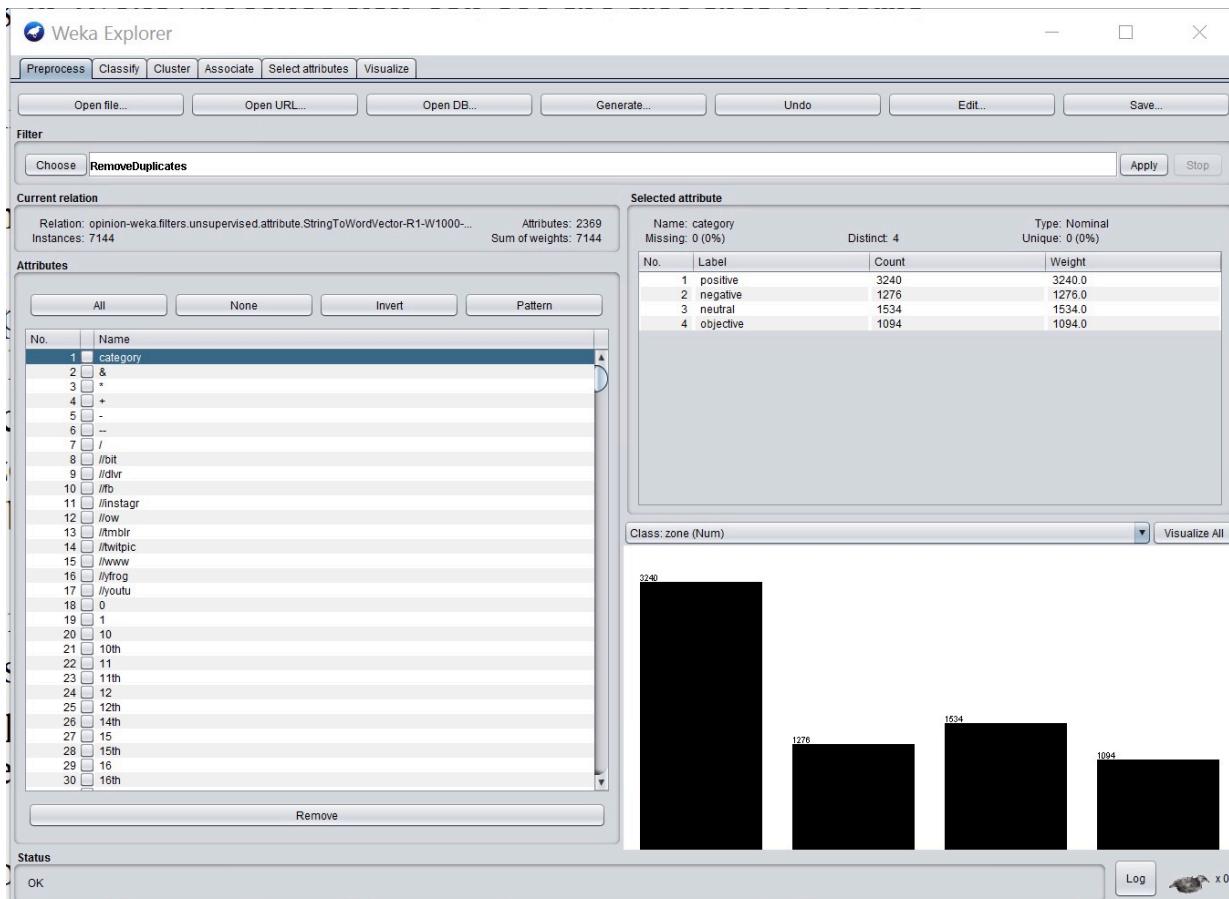


Figure 1.3 SMO

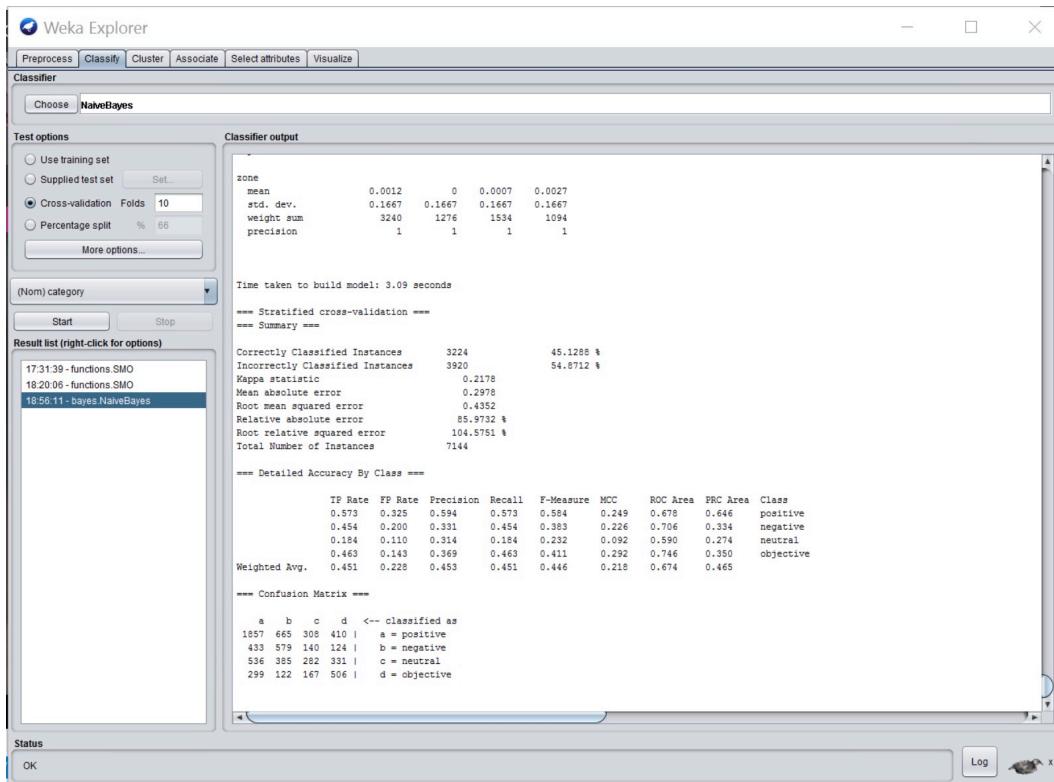
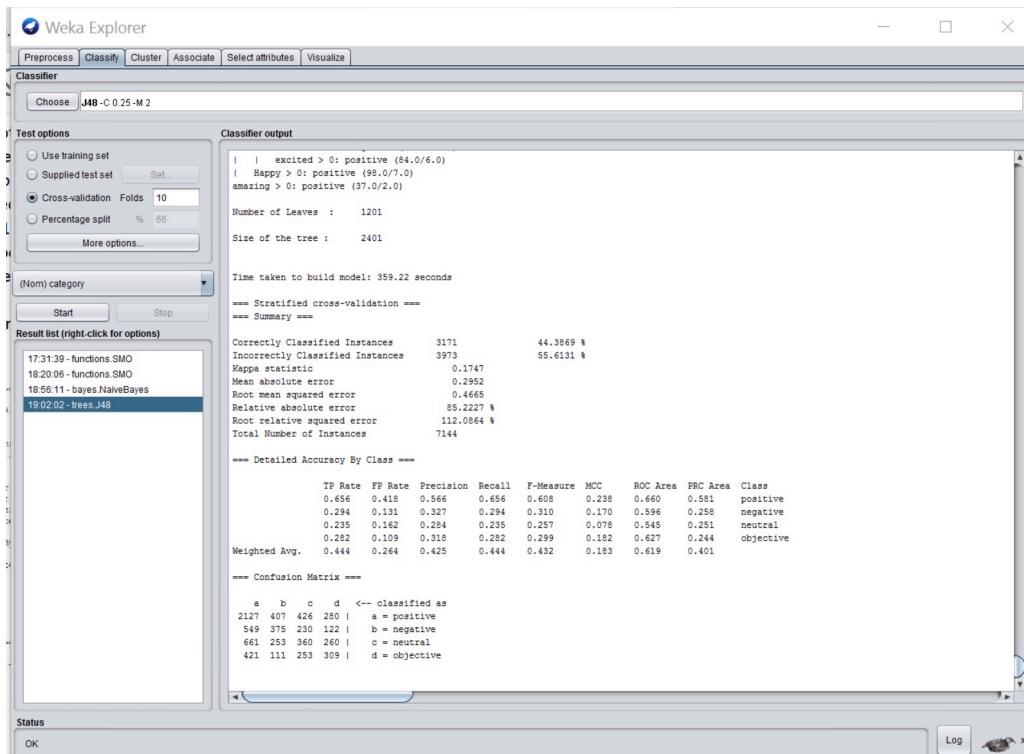


Figure 1.4 NaiveBay



In this part, we did not add any features to train classifiers. Instead, we simply used bag of words representation. That is we used all words in the message set as attributes to train classifiers including SVM, J48, and Naive Bayes.

- **First Experiment**

Input data file for Weka is the provided twitter file in .arff format: 'semeval_twitter_data.arff'. We kept all words originally in the file, and used StringToWordVector and RemoveDuplicate in Weka to extract words. The overall processing time was longer than expected since dimension of the vector space is high.

For SMO, we have 51.5258% accuracy.

For J48, we have 44.3869% accuracy.

For NB, we have 45.1288% accuracy.

Well, the accuracy of SMO is not so bad. The result is acceptable if not considering the processing time.

- **Second Experiment**

We decided to reduce the dimension of the vector space by using our Python script 'generate.py' to eliminate stop words, rare words, punctuation, emoji, and all special characters. The stop word list is located in folder 'resource'.

The outcome was not improved nor declined very significantly. The accuracy was about the same and the processing time was still long.

Part 2

To improve the classification result (accuracy), we did 5 different experiments, which each of them has a distinct set of features for the ML algorithms to train the classifiers. The result was getting better and better and the processing time was reduced significantly.

- **Frist Experiment**

Resource used: Positive word list (Positive.csv), Negative word list (Negative.csv)
Python Script: generate.py
Output folder is 'E1_pos_neg_result'.

Our script counts the number of positive&negative words in each twitter message, then writes data to a new .arff file, which is included in the output folder. The detailed prediction result of three classifiers is included in the output folder.

NB

```
== Stratified cross-validation ==
== Summary ==

Correctly Classified Instances      3180          43.9834 %
Incorrectly Classified Instances   4050          56.0166 %
Kappa statistic                   0.1099
Mean absolute error               0.3326
Root mean squared error          0.4151
Relative absolute error           96.042 %
Root relative squared error     99.7529 %
Total Number of Instances        7230

== Detailed Accuracy By Class ==

      TP Rate  FP Rate  Precision  Recall  F-Measure  MCC    ROC Area  PRC Area  Class
      0.759    0.563    0.529     0.759    0.623     0.205   0.634    0.586    positive
      0.000    0.000    ?         0.000    ?         ?       0.527    0.187    negative
      0.444    0.321    0.273     0.444    0.338     0.106   0.572    0.248    neutral
      0.000    0.000    ?         0.000    ?         ?       0.607    0.192    objective
Weighted Avg.      0.440    0.325    ?         0.440    ?         ?       0.597    0.382

== Confusion Matrix ==

  a   b   c   d  <- classified as
2493  0  791  0 |  a = positive
  790  0  503  0 |  b = negative
  860  0  687  0 |  c = neutral
  573  0  533  0 |  d = objective
```

J48

```
Time taken to build model: 0.09 seconds

== Stratified cross-validation ==
== Summary ==

Correctly Classified Instances      3284          45.4219 %
Incorrectly Classified Instances   3946          54.5781 %
Kappa statistic                   0
Mean absolute error               0.3359
Root mean squared error          0.4279
Relative absolute error           96.9977 %
Root relative squared error     102.8399 %
Total Number of Instances        7230

== Detailed Accuracy By Class ==

      TP Rate  FP Rate  Precision  Recall  F-Measure  MCC    ROC Area  PRC Area  Class
      1.000    1.000    0.454     1.000    0.625     ?     0.500    0.454    positive
      0.000    0.000    ?         0.000    ?         ?     0.500    0.179    negative
      0.000    0.000    ?         0.000    ?         ?     0.500    0.214    neutral
      0.000    0.000    ?         0.000    ?         ?     0.500    0.153    objective
Weighted Avg.      0.454    0.454    ?         0.454    ?         ?     0.500    0.307

== Confusion Matrix ==

  a   b   c   d  <- classified as
3284  0   0   0 |  a = positive
1293  0   0   0 |  b = negative
1547  0   0   0 |  c = neutral
1106  0   0   0 |  d = objective
```

SMO

```
==== Stratified cross-validation ====
==== Summary ====
Correctly Classified Instances      3284          45.4219 %
Incorrectly Classified Instances   3946          54.5781 %
Kappa statistic                      0
Mean absolute error                 0.3463
Root mean squared error              0.4161
Relative absolute error             99.9949 %
Root relative squared error        100          %
Total Number of Instances           7230

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC    ROC Area  PRC Area  Class
      1.000    1.000    0.454     1.000    0.625     ?     0.499    0.454    positive
      0.000    0.000    ?         0.000    ?         ?     0.499    0.179    negative
      0.000    0.000    ?         0.000    ?         ?     0.499    0.214    neutral
      0.000    0.000    ?         0.000    ?         ?     0.499    0.153    objective
Weighted Avg.       0.454    0.454    ?         0.454    ?         ?     0.499    0.307

==== Confusion Matrix ====

      a      b      c      d  <-- classified as
3284  0      0      0      |      a = positive
1293  0      0      0      |      b = negative
1547  0      0      0      |      c = neutral
1106  0      0      0      |      d = objective
```

• Second Experiment

Resource used: Tweepy and TextBlob python package

The API reference can be found at:

http://textblob.readthedocs.io/en/dev/api_reference.html

Python Script: generate.py, sentiment.py

Output folder is 'E2_API_polarity_result'.

The result from experiment 1 was not acceptable. Though processing time was short, we had a lower accuracy. The script sentiment.py imports two well-encapsulated python packages and it calculates a sentiment polarity score of each message. For this test, we only had the polarity scores as feature.

```

Number of Leaves : 14
Size of the tree : 27

Time taken to build model: 0.03 seconds

== Stratified cross-validation ==
== Summary ==

Correctly Classified Instances      3438          47.5519 %
Incorrectly Classified Instances   3792          52.4481 %
Kappa statistic                      0.0985
Mean absolute error                  0.3265
Root mean squared error              0.4044
Relative absolute error              94.2781 %
Root relative squared error         97.1869 %
Total Number of Instances           7230

== Detailed Accuracy By Class ==

      TP Rate  FP Rate  Precision  Recall  F-Measure  MCC  ROC Area  PRC Area  Class
0.929    0.821    0.485     0.929    0.637    0.161   0.670    0.589    positive
0.283    0.087    0.416     0.283    0.337    0.230   0.658    0.299    negative
0.000    0.001    0.000     0.000    0.000    -0.012   0.565    0.247    neutral
0.019    0.006    0.368     0.019    0.036    0.053   0.615    0.208    objective
Weighted Avg.    0.476    0.389    0.351     0.476    0.355    0.120   0.637    0.405

== Confusion Matrix ==

a   b   c   d  <-- classified as
3051 215  3  15 |  a = positive
 919 366  1   7 |  b = negative
1340 193  0  14 |  c = neutral
 979 106  0   21 |  d = objective

```

SMO

```

== Stratified cross-validation ==
== Summary ==

Correctly Classified Instances      3284          45.4219 %
Incorrectly Classified Instances   3946          54.5781 %
Kappa statistic                      0
Mean absolute error                  0.3349
Root mean squared error              0.4268
Relative absolute error              96.728 %
Root relative squared error         102.5774 %
Total Number of Instances           7230

== Detailed Accuracy By Class ==

      TP Rate  FP Rate  Precision  Recall  F-Measure  MCC  ROC Area  PRC Area  Class
1.000    1.000    0.454     1.000    0.625    ?    0.500    0.454    positive
0.000    0.000    ?        0.000    ?        ?    0.539    0.209    negative
0.000    0.000    ?        0.000    ?        ?    0.503    0.215    neutral
0.000    0.000    ?        0.000    ?        ?    0.500    0.153    objective
Weighted Avg.    0.454    0.454    ?        0.454    ?        ?    0.508    0.313

== Confusion Matrix ==

a   b   c   d  <-- classified as
3284  0   0   0 |  a = positive
1293  0   0   0 |  b = negative
1547  0   0   0 |  c = neutral
1106  0   0   0 |  d = objective

```

NB

```
==== Stratified cross-validation ====
==== Summary ====
```

Correctly Classified Instances	3284	45.4219 %
Incorrectly Classified Instances	3946	54.5781 %
Kappa statistic	0	
Mean absolute error	0.3463	
Root mean squared error	0.4161	
Relative absolute error	100	%
Root relative squared error	100	%
Total Number of Instances	7230	

```
==== Detailed Accuracy By Class ====
```

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
1.000	1.000	0.454	1.000	0.625	?	?	0.499	0.454	positive
0.000	0.000	?	0.000	?	?	?	0.499	0.179	negative
0.000	0.000	?	0.000	?	?	?	0.499	0.214	neutral
0.000	0.000	?	0.000	?	?	?	0.499	0.153	objective
Weighted Avg.	0.454	0.454	?	0.454	?	?	0.499	0.307	

```
==== Confusion Matrix ====
```

a	b	c	d	<-- classified as
3284	0	0	0	a = positive
1293	0	0	0	b = negative
1547	0	0	0	c = neutral
1106	0	0	0	d = objective

• Third Experiment

On top of the first and second experiment .arff data files, we added and combined the features. New features includes a weight of a message; it was calculated by generate.py. Based on 'TaboadaGrieve2004-SO.csv', we accumulated each word's weight; the more positive a message is the higher weight it gets. Overall, we had positive words count, negative words count, weight, polarity, and if contains question mark as features. We had a higher accuracy of J48.

J48

```
Number of Leaves : 221
```

```
Size of the tree : 441
```

```
Time taken to build model: 0.11 seconds
```

```
==== Stratified cross-validation ====
==== Summary ====
```

Correctly Classified Instances	3500	48.4094 %
Incorrectly Classified Instances	3730	51.5906 %
Kappa statistic	0.1578	
Mean absolute error	0.3112	
Root mean squared error	0.407	
Relative absolute error	89.883 %	
Root relative squared error	97.8275 %	
Total Number of Instances	7230	

```
==== Detailed Accuracy By Class ====
```

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
0.854	0.663	0.518	0.854	0.645	0.220	0.701	0.623	positive	
0.252	0.061	0.475	0.252	0.329	0.250	0.647	0.318	negative	
0.154	0.085	0.331	0.154	0.211	0.095	0.595	0.268	neutral	
0.117	0.044	0.322	0.117	0.171	0.114	0.657	0.245	objective	
Weighted Avg.	0.484	0.337	0.440	0.484	0.423	0.183	0.662	0.435	

```
==== Confusion Matrix ====
```

a	b	c	d	<-- classified as
2806	153	209	116	a = positive
769	326	141	57	b = negative
1078	131	239	99	c = neutral
769	76	132	129	d = objective

SMO

```
==== Stratified cross-validation ====
==== Summary ====
Correctly Classified Instances      3284          45.4219 %
Incorrectly Classified Instances   3946          54.5781 %
Kappa statistic                      0
Mean absolute error                 0.333
Root mean squared error              0.4245
Relative absolute error              96.1555 %
Root relative squared error         102.0177 %
Total Number of Instances           7230

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC    ROC Area  PRC Area  Class
      1.000    1.000    0.454     1.000    0.625     ?     0.500    0.454    positive
      0.000    0.000    ?         0.000    ?         ?     0.607    0.241    negative
      0.000    0.000    ?         0.000    ?         ?     0.505    0.216    neutral
      0.000    0.000    ?         0.000    ?         ?     0.520    0.159    objective
Weighted Avg.      0.454    0.454    ?         0.454    ?         ?     0.523    0.320

==== Confusion Matrix ====

      a      b      c      d  <-- classified as
  3284  0      0      0 |  a = positive
  1293  0      0      0 |  b = negative
  1547  0      0      0 |  c = neutral
  1106  0      0      0 |  d = objective
```

NB

```
==== Stratified cross-validation ====
==== Summary ====
Correctly Classified Instances      3233          44.7165 %
Incorrectly Classified Instances   3997          55.2835 %
Kappa statistic                      0.2115
Mean absolute error                 0.3055
Root mean squared error              0.4084
Relative absolute error              88.2162 %
Root relative squared error         98.1463 %
Total Number of Instances           7230

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC    ROC Area  PRC Area  Class
      0.640    0.302    0.638     0.640    0.639    0.337  0.715    0.685    positive
      0.187    0.042    0.490     0.187    0.271    0.220  0.666    0.345    negative
      0.165    0.099    0.313     0.165    0.217    0.086  0.623    0.288    neutral
      0.573    0.325    0.242     0.573    0.340    0.186  0.674    0.237    objective
Weighted Avg.      0.447    0.216    0.481     0.447    0.437    0.239  0.680    0.470

==== Confusion Matrix ====

      a      b      c      d  <-- classified as
  2101  109   254   820 |  a = positive
  398   242   167   486 |  b = negative
  513   93    256   685 |  c = neutral
  282   50    140   634 |  d = objective
```

- **Fourth Experiment**

Since the highest accuracy so far has not reached beyond 50%. We were thinking the reasons that restrains the accuracy. Reasons we concluded:

- The features we were using were not strong enough to classify a message to one of the four categories. Since so far we only have features tell us if a message tends to approach positive or negative. Features, such as positive/negative word number, weight, and polarity score, only give information of which direction a message might be approaching to. However, we have four directions that a message can approach. Thus, the vector space might be skewed.

This time we tried a new set of features includes the positive sentiment score, neutral score, negative score, and a compound score. We made use of a existing well encapsulated package ‘nltk.sentiment’ package. The classification results of three classifiers has improved and reached beyond 50 percent.

J48

```

Number of Leaves : 11
Size of the tree : 21

Time taken to build model: 0.04 seconds

== Stratified cross-validation ==
== Summary ==

Correctly Classified Instances      3815          52.7663 %
Incorrectly Classified Instances   3415          47.2337 %
Kappa statistic                      0.309
Mean absolute error                  0.295
Root mean squared error              0.3882
Relative absolute error              85.1848 %
Root relative squared error        93.2879 %
Total Number of Instances           7230

== Detailed Accuracy By Class ==

      TP Rate  FP Rate  Precision  Recall  F-Measure  MCC  ROC Area  PRC Area  Class
      0.717    0.255    0.701    0.717    0.709    0.461  0.738   0.654   positive
      0.542    0.129    0.478    0.542    0.508    0.394  0.735   0.404   negative
      0.450    0.259    0.321    0.450    0.375    0.171  0.617   0.282   neutral
      0.058    0.029    0.268    0.058    0.095    0.059  0.696   0.264   objective
Weighted Avg.      0.528    0.199    0.514    0.528    0.508    0.326  0.705   0.470

== Confusion Matrix ==

     a      b      c      d  <-- classified as
2354  333  531  66 |   a = positive
 293  701  273  26 |   b = negative
 463  305  696  83 |   c = neutral
 250  126  666  64 |   d = objective

```

SMO

==== Stratified cross-validation ====
==== Summary ===

Correctly Classified Instances	3820	52.8354 %
Incorrectly Classified Instances	3410	47.1646 %
Kappa statistic	0.2915	
Mean absolute error	0.3183	
Root mean squared error	0.4065	
Relative absolute error	91.9149 %	
Root relative squared error	97.6952 %	
Total Number of Instances	7230	

==== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
0.755	0.332	0.655	0.755	0.701	0.422	0.721	0.621	positive	
0.462	0.093	0.520	0.462	0.489	0.387	0.729	0.381	negative	
0.480	0.273	0.324	0.480	0.387	0.183	0.609	0.271	neutral	
0.000	0.000	?	0.000	?	?	0.637	0.207	objective	
Weighted Avg.	0.528	0.226	?	0.528	?	?	0.686	0.440	

==== Confusion Matrix ===

a	b	c	d	<-- classified as
2481	243	560	0	a = positive
405	597	291	0	b = negative
589	216	742	0	c = neutral
315	93	698	0	d = objective

NB

==== Stratified cross-validation ====
==== Summary ===

Correctly Classified Instances	3779	52.2683 %
Incorrectly Classified Instances	3451	47.7317 %
Kappa statistic	0.3164	
Mean absolute error	0.2602	
Root mean squared error	0.4059	
Relative absolute error	75.1524 %	
Root relative squared error	97.5478 %	
Total Number of Instances	7230	

==== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
0.722	0.266	0.693	0.722	0.708	0.455	0.770	0.737	positive	
0.470	0.098	0.510	0.470	0.489	0.384	0.769	0.455	negative	
0.047	0.030	0.303	0.047	0.082	0.040	0.654	0.297	neutral	
0.656	0.269	0.306	0.656	0.417	0.297	0.735	0.285	objective	
Weighted Avg.	0.523	0.186	0.518	0.523	0.490	0.329	0.740	0.523	

==== Confusion Matrix ===

a	b	c	d	<-- classified as
2372	269	46	597	a = positive
310	608	92	283	b = negative
480	224	73	770	c = neutral
259	91	30	726	d = objective

- **Fifth Experiment (Best accuracy achieved)**

Resource used: NLTK python package

The API reference can be found at:

http://www.nltk.org/api/nltk.sentiment.html#module-nltk.sentiment.sentiment_analyzer

Python Script: generate.py, NLTKAPI.py

Output folder is 'E5_polarity_sentimnet_result'.

There is no new features included. We used polarity score and compound score calculated by nltk.sentiment package as features. We got the highest accuracy (53.029%).

J48

```

Number of Leaves : 43
Size of the tree : 85

Time taken to build model: 0.04 seconds

==== Stratified cross-validation ====
==== Summary ====

    Correctly Classified Instances      3834           53.029 %
    Incorrectly Classified Instances   3396           46.971 %
    Kappa statistic                   0.3092
    Mean absolute error              0.294
    Root mean squared error          0.3876
    Relative absolute error          84.9048 %
    Root relative squared error     93.1495 %
    Total Number of Instances       7230

==== Detailed Accuracy By Class ====


|               | TP Rate | FP Rate | Precision | Recall | F-Measure | MCC   | ROC Area | PRC Area | Class     |
|---------------|---------|---------|-----------|--------|-----------|-------|----------|----------|-----------|
| positive      | 0.760   | 0.322   | 0.662     | 0.760  | 0.708     | 0.436 | 0.742    | 0.658    | positive  |
| negative      | 0.513   | 0.122   | 0.478     | 0.513  | 0.495     | 0.380 | 0.737    | 0.418    | negative  |
| neutral       | 0.131   | 0.072   | 0.332     | 0.131  | 0.188     | 0.087 | 0.624    | 0.298    | neutral   |
| objective     | 0.427   | 0.162   | 0.323     | 0.427  | 0.367     | 0.237 | 0.695    | 0.273    | objective |
| Weighted Avg. | 0.530   | 0.208   | 0.507     | 0.530  | 0.506     | 0.321 | 0.709    | 0.479    |           |


==== Confusion Matrix ====


| a    | b   | c   | d   | <-- classified as |
|------|-----|-----|-----|-------------------|
| 2496 | 306 | 150 | 332 | a = positive      |
| 367  | 663 | 96  | 167 | b = negative      |
| 565  | 287 | 203 | 492 | c = neutral       |
| 340  | 131 | 163 | 472 | d = objective     |


```

SMO

```
== Stratified cross-validation ==
== Summary ==

Correctly Classified Instances      3652          50.5118 %
Incorrectly Classified Instances   3578          49.4882 %
Kappa statistic                      0.1638
Mean absolute error                  0.3283
Root mean squared error              0.4187
Relative absolute error              94.7942 %
Root relative squared error        100.6198 %
Total Number of Instances           7230

== Detailed Accuracy By Class ==

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC    ROC Area  PRC Area  Class
0.928     0.753    0.506     0.928    0.655     0.233   0.591    0.506    positive
0.464     0.099    0.506     0.464    0.484     0.378   0.703    0.356    negative
0.004     0.004    0.214     0.004    0.008     0.000   0.496    0.213    neutral
0.000     0.000    ?         0.000    ?         ?       0.372    0.133    objective
Weighted Avg.  0.505     0.360    ?         0.505    ?         ?       0.557    0.359

== Confusion Matrix ==

      a     b     c     d  <-- classified as
3046  227   11   0 |   a = positive
 686  600   7   0 |   b = negative
1289  252   6   0 |   c = neutral
 996  106   4   0 |   d = objective
```

NB

```
== Stratified cross-validation ==
== Summary ==

Correctly Classified Instances      3783          52.3237 %
Incorrectly Classified Instances   3447          47.6763 %
Kappa statistic                      0.2924
Mean absolute error                  0.2826
Root mean squared error              0.3856
Relative absolute error              81.6015 %
Root relative squared error        92.6711 %
Total Number of Instances           7230

== Detailed Accuracy By Class ==

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC    ROC Area  PRC Area  Class
0.770     0.339    0.654     0.770    0.707     0.430   0.779    0.746    positive
0.440     0.099    0.491     0.440    0.464     0.356   0.755    0.462    negative
0.191     0.115    0.312     0.191    0.237     0.093   0.652    0.294    neutral
0.353     0.142    0.310     0.353    0.330     0.200   0.717    0.275    objective
Weighted Avg.  0.523     0.218    0.499     0.523    0.505     0.310   0.738    0.526

== Confusion Matrix ==

      a     b     c     d  <-- classified as
2528  229   235   292 |   a = positive
 366  569   184   174 |   b = negative
 608  241   296   402 |   c = neutral
 364  119   233   390 |   d = objective
```

- **Experiment out of Curiosity**

The highest accuracy so far is about 53%. According to statistics probability, the probability that we correctly classify a message to one of the four class labels is 1/4. Therefore, our accuracy is acceptable. We were curious why it could not reach a higher result. First guess would be our calculated sentiment score did not specially consider the ‘objective’ category. To prove our assumption, we did a sixth experiment.

In this test, when parsing all the messages, we changed ‘objective’ label to ‘neutral’. We got a much higher accuracy — 63.444%.

J48 pruned tree

```
compund_score <= 0.3695
| neu_score <= 0.844
| | compund_score <= 0.1551: negative (1158.0/564.0)
| | compund_score > 0.1551: positive (258.0/119.0)
| neu_score > 0.844: neutral (2759.0/1072.0)
compund_score > 0.3695: positive (3055.0/839.0)
```

Number of Leaves : 4

Size of the tree : 7

Time taken to build model: 0.03 seconds

==== Stratified cross-validation ====

==== Summary ===

Correctly Classified Instances	4587	63.444 %
Incorrectly Classified Instances	2643	36.556 %
Kappa statistic	0.4149	
Mean absolute error	0.253	
Root mean squared error	0.3571	
Relative absolute error	80.6825 %	
Root relative squared error	90.1975 %	
Total Number of Instances	7230	

SMO

==== Stratified cross-validation ====

==== Summary ====

Correctly Classified Instances	4559	63.0567 %
Incorrectly Classified Instances	2671	36.9433 %
Kappa statistic	0.4025	
Mean absolute error	0.2909	
Root mean squared error	0.3716	
Relative absolute error	92.7736 %	
Root relative squared error	93.8499 %	
Total Number of Instances	7230	

The last experiment was more biased, since we simply changed all ‘neutral’ label to ‘positive’.

J48

==== Stratified cross-validation ====

==== Summary ====

Correctly Classified Instances	6005	83.0567 %
Incorrectly Classified Instances	1225	16.9433 %
Kappa statistic	0.3186	
Mean absolute error	0.1231	
Root mean squared error	0.2485	
Relative absolute error	83.718 %	
Root relative squared error	91.7041 %	
Total Number of Instances	7230	